

# Operating Systems 2

13303

## Introduction

- Chapter 1:
  - Process creation
  - process termination
  - process scheduling
  - process intercommunication (IPC):
    - IPC with pipes
    - IPC with signals:
    - IPC with shared memory:
- Chapter 2:
  - memory management
    - swapping
    - virtual memory
    - segmentation
    - page replacement algorithms
    - Etc...
- Chapter 3:
  - File management system:
    - file system
      - user view
      - system view
      - Unix file system structure
- Chapter 4:
  - I/O device management
- Refresh from OS1
  - Process: program in execution
  - Program: set of instructions
  - Process state:
    - Ready state: process in ready queue (in memory), waiting for the scheduler to choose it
    - Running state: being executed (occupies the CPU)
    - Waiting/Blocking state: I/O request
  - Timings:
    - 1: New → Ready: process is put into ready queue
    - 2: Ready → Running: Scheduler keeps the CPU to the process
    - 3: Running → Ready: The scheduler picks up the process and gives the CPU to another process

- 4: Running → Blocked: Requests for I/O
- 5: Blocked → Ready: I/O are arriving
- 6: Running → End

- **Process Implementation:**

- Divided into 4 parts:
  - Code: instructions to execute
    - instructions: I1 → I2 → I3
    - PC: program counter (contains the address of the next instruction to execute)
  - Data: Global & static variables
  - Stack: Contains function calls & local variables
  - Heap: Dynamic allocation of memory
- To implement a process, the OS creates a table in memory
  - Table of processes: array of structures where each row is a structure
  - This table & process is the PCB: process control block
- What does the PCB include:
  - pid: id of the process
  - state
  - priority
  - PC value
  - PSW: process status word
    - 2 modes of execution: User mode or Kernel mode
    - it is a bit that saves the mode of execution
  - pointer to code
  - pointer to data
  - stack
  - real user id
  - parent id
  - root directory
  - exit status
  - etc...
- Context Switching: saving the state of the process (PCB ig)
- CPU registers (not all):
  - IR: instruction register (Fetch)
  - PC
  - PSW
- Command to get the programs in execution: ps

- **Process Creation:**

- By duplication, except the first process of *pid=1* named *init*
- Example:
  - P1 process with stack, heap, data, and code
  - P2 is created by duplicating P1's stack, heap, and data except the code
- How to duplicate:
  - [library <unistd.h>](#) containing a function

- function: `int fork()`
- the child starts running the program starting from the fork it was created (exclusive)
- Return value of the fork informs the current pid of the process
  - if return is -1: failure → no child created
  - if return is 0: I am in the child
  - else it returns the pid of the child and I am in the parent

○ Example:

- Parent: pid, variables, file descriptor
- child: pid, copy of variables, another file descriptor that points to the same file
- However: the position inside the file is shared in both parent & child

**11/30/2022**

• **Process Identification:**

- `int getpid()` → returns the id of the calling process
- `int getppid()` → returns the id of the parent

• **Process termination:**

○ Normal exit:

- finish the execution of all instructions
- call `exit()`
  - **`void exit(int status)`**
    - library: `<sys/wait.h>`
    - the process calling exit ends its execution
    - status
      - used by the process to inform his parent about the exit status. it is between 0 & 255 (1 byte)
      - value is stored on the left of this byte. So to actually access the real data, we need to shift the status to the right by 8 bits (1 byte)
    - By convention, `exit(0)` is a normal exit
  - Parent calls `fork()` → a child is created
    - The parent & the child continue execution
    - the child decided to `exit(status)`, but only exists if the parent knows that the child exited
    - To keep track of the child, we use a function in the parent:
      - `wait(&b)`
      - `b>>8` → shifts the status 8 bits in order to read it

○ Abnormal:

- It is killed by a privileged process

○ **`int wait(int *st):`**

- function called by the parent to wait for the exit of **one** child
- this statement is blocking (the parent stops execution until the child exits)
- it returns the pid of the exited child

○ **`int waitpid(int pid, int *st, int options):`**

- waits for a specific child with **pid** to exit

- `waitpid(-1, ... , ...)` → `wait()`
- **options:**
  - 0 → nothing
  - `WNOHANG` → makes the wait a non-blocking statement
- Macros:
  - **`int WIFEXITED(int status)`**: returns non-zero value if the child terminated normally with exit
    - Example:
      - `wait(&str)`

```

if WIFEXITED(str)
{
    printf("my child exited normally");
}

```
  - **`int WEXITSTATUS(int status)`**: shifts and returns the status automatically
    - Example:
      - `wait(&str)`

```

if WIFEXITED(str)
{
    printf("my child exited normally with value %d", WEXITSTATUS(str) );
}

```
  - **`int WIFSIGNALED(int status)`**
    - returns a non-zero value if the child terminated abnormally because it received a signal that was not handled
  - **`int WTERMSIG(int status)`**
    - informs us of the signal that killed the child
  - **`int sleep(int sec)`**:
    - the process suspends its execution sec time (blocking statement)
- If the **parent dies before its child** → child is called an **Orphan process**
  - The init adopts this child & the ppid of this child becomes 1
- If the **child dies and the parent doesn't execute the wait** statement → the child is a **zombie process**
- Command:
  - **`ps-a pid`** → this command gives the status of the child with **pid**.
    - if orphan → O
    - if zombie → Z
    - if normal → 0
- **Exercise:** Write a program that determines the maximum of a vector of integers by distributing the task between a parent process and his child.

**5/12/2022**

## Executing a file

- **Shell:** a process executed by the system to write commands
- Inside the shell I am writing a command (example `ls -l -a`):
  - the shell needs to execute its own instructions? How will it execute the `ls` command?
- **Process:**
  - shell executes `fork()` → child has a `pid = 101`
  - I make the child execute the command `ls`
- **NOTE:**
  - `void main(int argc, char*argv[])`
    - Command line arguments: arguments of the main function. Written during the call to execution
    - Example: `ls -l -a`
    - `argc`: number of parameters (including the name of the program)
      - In the example: `argc = 3`
    - `argv`: An array of strings containing the parameter values
      - in the example: `argv = [ "ls", "-l", "-a", "\0" ]`
  - What is a path:
    - a variable, indicating the path to reach a specific folder or file
- Execution of the "`ls -l -a`" command:
  - 1) `main (int argc, char*argv[])`
  - 2) parent executes `fork()`
  - 3) inform the child to execute a new image/process ( in this case: `ls -l -a`)
    - The child no longer is a shell, it is a program executing the command
- How to change the main task of the child to another task?
- **Family of exec functions**
  - Role: replace the current context of a process with new image/process
  - These functions don't return
    - they never return to continue the instructions under it
  - Functions:
    - `int execl (char*filename, char* arg0, char* arg1,...)`
      - `execl (list)`
      - filename : should be the full path
    - `int execlp (char*filename, char* arg0,char*arg1,...)`
      - filename: only the file name (no need for a path)
    - `int execv (char*filename, char*argv[])`
      - `execv (vector)`
      - filename: should be a full path
    - `int execlp (char*filename, char*argv[])`
      - filename: only the file name (no need for a path)
- Example: write a program to execute the command "`ls -l -a`"

```

void main(int argc, char* argv[])
{
    int pid;
    pid = fork();
    if(!pid)
    {
        execl("/bin/ls","ls","-l","-a", NULL); //filename is a path
        //execlp("ls","ls","-l","-a",NULL); //filename is just a name
        printf("Never executed\n"); //this process is never executed, since the child
                                   never returns to this process after exec functions
    }
}



```

- Return value of the exec:
  - -1 → failure
  - 0 → success
  - This return value is returned to the parent

**7/12/2022**


### **Fork & Binary Tree**

- A && B
  - are 2 boolean expressions
  - if A is false, then B is never executed/evaluated
  - fork && fork:
    - Both forks are done in the parent
- A || B:
  - if A is true → B is never executed
  - fork || fork:
    - 1<sup>st</sup> is done in the parent, 2<sup>nd</sup> is done in the child



### **Interprocess Communication (IPC)**

- Methods:
  - Naïve: Wait & exit
    - very limited/inefficient
  - Unix pipes
  - Signals
  - Shared memory segment
  - Message passing
  - Communication by sockets (when processes are on different machines)
- Unix Pipes:
  - ls -l | wc
    - the output of ls -l, is the input of the wc process
  - Definition: A pipe is an [interprocess communication tool](#). It is a [FIFO](#) buffer (like a queue)
    - **buffer**: zone in memory for the OS
    - It is a [one way communication tool](#).
    - Size of the pipe is **fixed** in the OS (hyper-parameter)



- usually it is 4KB
- The pipe has 2 sides for access
  - **By convention: one side is for writing, & the other side is for reading**
- Pipe has 2 tasks: communication & synchronization

#### ○ Pipe creation:

- we need an **array of 2 integers**
  - `int fd[2];`
- **`int pipe(fd);`** → function that creates a pipe
  - `fd[0]` is the reading side
  - `fd[1]` is the writing side
  - return value:
    - 0 → success
    - -1 → fail (wasn't created)

#### ○ I/O with pipes:

- In UNIX: Every device is a file
  - Everything is represented as a file that can be opened, read, and edited through a descriptor (just like a file pointer)
- `fd[0]` stores the descriptor for **reading**, `fd[1]` stores the descriptor for **writing**
- **`int write(fd[1], buf, size)`**
  - size: number of characters to write
  - buffer: the zone to write the data
  - remove from the buffer size characters and put them in `fd[1]`
- **`int read(fd[0], buf, size):`**
  - size: number of characters to read
  - buffer: the zone to read the data
  - read from `fd[0]` size characters and put them in buffer

```
int x[10];
read(fd[0], &x, 10 * sizeof(int));
```

- Both return the number of effective characters

#### ○ Synchronization with pipes:

- Unix pipes are used as tools of synchronization between processes, How?
- **Reading from empty pipe is a blocking statement if there is a writer in other side**
  - pipe is empty while the writer is still not closed. So the pipe waits until the writer writes something in order to read it
- **Writing in full pipe is a blocking statement if there is a reader in other side**
  - pipe is full while the reader is not closed. pipe waits until the reader reads something so it gives the pipe space to write

#### ○ Pipe & fork()


- If I want to use pipe as a way of communication, I need to create a pipe before forking (so I can communicate between the parent and the child)

- **How** to connect 2 processes using pipes:
  - 1) Make the pipe using the pipe() function
  - 2) fork() to create the child for reading from pipe
  - 3) **in the child, close the writing side (if child wants to read)**
  - 4) **in the parent, close the reading side (if parent wants to write)**

```
void main()
{
    int fd[2], pid;
    pipe(fd); //create the pipe --> parent has a pipe with fd[1] (write) & fd[0](read)

    int p = fopen("toto.txt");
    parent --> close(fd[0]); //if the parent wants to write
    child --> close(fd[1]); //if the child wants to read

    /*This is one way of communication, If I want 2 ways of communication between parent &
    child, I have to create another pipe with opposite closing channels*/
}
```



- **Descriptors:**
  - By default, each created process is associated with three file descriptors:
    - 0 → standard input (Keyboard)
    - 1 → standard output (Screen)
    - 2 → standard error (Screen) (for error messages)
  - In Pipes, it is associated with 2 descriptors:
    - fd[0] for reading
    - fd[1] for writing
- **In forking;**
  - the pipe is a **shared zone between parent & child**
  - However, the **child & parent have different descriptors, each pointing to the same location in the file**
    - child fd[0] != parent fd[0]
    - but child fd[0] points to the same place in the parent fd[0]

- **Example:**

```
void main()
{
```



```

char reception[100]; //for receiving the message
char *message = "Hello my dear";
int desc[2],nb;

pipe(desc);
if(!fork())
{
    close(desc[1]);
    nb = read(desc[0],reception,100);
    printf("the child read %d characters --> %s\n",nb,reception);
    close(desc[0]);
}
else{
    close(desc[0]);
    write(desc[1],message,strlen(message)+1);
    close(desc[1]);
}
}

```

**12/12/2022**

- Steps to execute: ls -l | wc

- Make a pipe
- Create a child
- I want the parent to execute ls -l → parent writes → close fd[0]
- I want the child to execute wc → child reads → close fd[1]

```

void main(){
    int fd[2];
    pipe(fd);

    if(fork())
    {
        //in parent
        //Missing: how can i send the output to the child?
        execlp("ls","ls","-l",NULL);
    }
    else{
        //in child
        //Missing: how can i read the output?
        execlp("wc","wc",NULL);
    }
}

```

- Redirection of I/O:

- Each process has by default an associated table called “file descriptor table”.
- File descriptor table contains rows, each containing pointers to the file objects which do all the resource handling.

- All file descriptors have, **by default**, the values:

0 (keyboard)
1 (screen)
2 (screen for error)

- Each time a **file opens, a descriptor is included in the table**

- `int p = open("toto.txt")`
- A file descriptor of index (for example) 100 is added to the file descriptor table

0 (keyboard)
1 (screen)
2 (screen for error)
100

- When a **pipe is created, 2 descriptors `fd[1]` and `fd[2]` are added** to the file descriptor table

0 (keyboard)
1 (screen)
2 (screen for error)
<code>fd[0]</code> (read)
<code>fd[1]</code> (write)

- **By default:**

- `ls -l` writes to descriptor 1,
- `wc` reads from descriptor 0.

- However, the only communication between the parent and the child is the pipe. So I **want to make `ls -l` write to `fd[1]` and `wc` to read from descriptor `fd[0]`**

- To execute `ls -l | wc`:

- I should redirect the output of the process "`ls`" from screen to the writing side of the pipe
- We should redirect the input of the process "`wc`" from the keyboard to the reading side of the pipe

- **How to redirect the I/O:**

- **The 2 duplicate functions:**

- `int dup(int newfd)`

- it function **duplicates the content of a file descriptor**, given as a parameter, **in the first free cell** of the file descriptor table
- So to put the fd[1] instead of 1 in the ls -l, I have to free 1.
- How to free a cell:
  - `Close(1);`
- Then I do: `dup(fd[1]);`
- Now there **exists 2 fd[1]** in the file descriptor table → `close(fd[1]);`
- **Int dup2(int newfd,int oldfd):**
  - **Automatically** does the first close for me
  - `Dup2(fd[1],1); == close(1);dup(fd[1]);`

```
void main(){
    int fd[2];
    pipe(fd);

    if(fork())
    {
        //in parent
        close(fd[0]);
        dup2(fd[1],0);
        close(fd[1]);
        execlp("ls","ls","-l",NULL);
    }
    else{
        //in child
        close(fd[1]);
        dup2(fd[0],0);
        close(fd[0]);
        execlp("wc","wc",NULL);
    }
}
```

- All the above where called **UNNAMED pipes**:
  - They act **like any variable** (deleted after program ends)
- **Named pipes**:
  - It's like a file, **has its own existence**. Other processes **can access it later on**
  - A named pipe or "FIFO Special file" is similar to pipe but instead of being an anonymous, temporary connection, **it has a name like any other file and permanent existence**.
  - The process opens the FIFO by its name in order to communicate through it
  - The named pipe has a **capacity of 40Kb**
  - The named pipe **can be used by independent processes other than the one that created it**
- **Creation of FIFO**:
  - `Int mkfifo(char *filename, mode_t mode)`
    - Mode can be: 0666, 0777
    - returns:

- -1 → failure
- 0 → success
- Create FIFO in terminal:
  - `mkfifo canal` → creates a named pipe named canal
- Using FIFO:
  - Open FIFO:
    - `Int fd = open("canal",...);`

**14/12/2022**

- Monday lab → bring a laptop
- Named pipe example (discussed in previous lecture):
  - writer.c

```
void main(){
    int n;
    char Buffer[100];

    int fd_write;

    mkfifo("pipe1",0777);

    fd_write = open("pipe1",O_WRONLY); //Open for writing only
    write(fd_write,"Bonjour",7);
    close(fd_write);
}
```

- reader.c

```
void main(){
    //in reader, i don't create the pipe since the pipe was created in writer.c
    int n;
    char buffer[100];
    int fd_read;
    fd_read = open("pipe1",O_RDONLY); //Open for reading only
    read(fd_read,buffer,100);
    buffer[100] = '\0';
    printf("%s\n",buffer);
    close(fd_read);
}
```

- Exercise: What would happen in the following programs? (Previous exam question)
  - Program 1:

```
void main(){
    int i,j=1,p[2];
```

```

pipe(p);
write(p[1], &j, sizeof(int));
for(i=1; i<5; i++){
    if(!fork()){
        close(p[1]);
        break;
    }
}
read(f[0], &j, sizeof(int));
printf("%d\n", i);
}

```

▪ Answer:

1) Pipe is created →



Parent file descriptor
0
1
2
p[0]
p[1]

2) The parent writes to the pipe the value of j=1.

3) For loop:

a. Creation of the child

b. If I am in the child → close its writing side p[1]



Child file descriptor
0
1
2
p[0]
<del>p[1]</del>

4) Here reading depends on the order of execution

a. If one of the children finished first:

i. Child reads 1 & removes it from the pipe & prints its i

ii. For all the others: Reading from an empty pipe but we have one of the writer sides open (parent fd[1] is still open) → blocking statement

b. If the parent P0 finishes first:

i. Parent reads 1 & removes it from the pipe & prints its i=5

ii. Parent ends its program → all its file descriptors are closed (including its fd[1])

iii. Children: reading from an empty pipe, but all the writer fd[1] are closed → not a blocking statement → all of them print their i values & exit.

○ Program 2:

```

void main(){
    int i,j=1,p[2];
    pipe(p);
    for(int i=1; i<5; i++){

```

```

        if(!fork()){
            close(p[1]);
            break;
        }
    }
    wait(NULL);
    read(p[0],&j,sizeof(int));
    printf("%d\n",i);
}

```

▪ Answer:


1) Pipe is created →

Parent file descriptor
0
1
2
p[0]
p[1]

2) For loop:

a. Creation of the child

b. If I am in the child → close its writing side p[1]



Child file descriptor
0
1
2
p[0]
<del>p[1]</del>

3) Parent waits for **one** of the children to finish

4) We go into one of the children:

a. Child reads from an empty pipe (we didn't write to the pipe) and p[1] of the parent open → blocking statement → no one exits

○ Program 3: **(HOMEWORK)**


```

void main(){
    int i,j=1,p[2];
    pipe(p);
    for(int i=0;i<5;i++){
        if(!fork()){
            close(p[1]);
            break;
        }
    }
    write(p[1],&j,sizeof(int));
    printf("%d\n",i);
    read(p[0],&j,sizeof(int));
}

```




**19/12/2022**

- Exercise 1:



```
void main(){
    int fd12[2],fd23[2],fd31[2];
    pipe(fd12);
    pipe(fd23);
    pipe(fd31);

    if(fork()) //I am in P1
    {
        close(fd31[1]);
        close(fd12[0]);
        close(fd23[1]);
        close(fd23[0]);
    }
    else if(fork()) //P2
    {
        close(fd12[1]);
        close(fd23[0]);
        close(fd31[0]);
        close(fd31[1]);
    }
    else{
        //P3
        close(fd12[0]);
        close(fd12[1]);
        close(fd23[1]);
        close(fd31[0]);
    }
}
```



- 2014 partial: Draw the graph


```
int main(){
```

```

if(fork())
{
    fork() && (fork() || (fork() && fork()));
    //A      B      C
}
else fork() || (fork() && fork());
//      D      E
}

```

○ Answer:



• Exercise:

- Parent creates 10 child processes
- Parent reads variable (N) from stdin and write 'm' N times in a pipe
- The 10 children start a race to read from the pipe
- Each child must send to the parent the number of characters 'm' he reads
- Finally, the parent must announce the winner (child with maximum m)
- Steps:
  - Parent:
    - Create the pipe
    - Read N
    - Create the 10 children
    - Write N 'm' in pipe
    - Determining the max
    - Display the winner
  - Child:
    - Start needing from pipe until no data exist
    - Send the counter to parent
    - Done

○ Answer:

```

int main(){

```



```

int Pid[10],status[10],N,parent_pid=getpid(),max,winner;
int fd[2];
char c = 'm';
pipe(fd);
close(fd[0]); //close reading side of father

//write N 'm' in pipe
for(int i=0;i<N;i++)
{
    write(fd[1],&c,sizeof(char));
}
//close writing side of parent
close(fd[1]);

//create 10 children
for(int i=0;i<10;i++)
{
    if(!(pid[i]=fork())){break;}
}


if(getpid() == parent_pid) //in parent
{
    for(int i=0;i<10;i++){waitpid(pid[i],&status[i],NULL);}

    max = WEXITSTATUS(status[0]);
    for(int i=0;i<10;i++)
    {
        if(max < WEXITSTATUS(status[i]))
        {
            max=WEXITSTATUS(status[i]);
            winner = pid[i];
        }
    }
}
else{ //in child
    sleep(2);
    close(fd[1]);
    while(read(fd[0],&c,sizeof(char)) != 0){count++;}
    exit(count);
}
}
//the scheduler handles the race between the pipes (round robin,FIFO...)
// I don't have to write in my code that another child has to come and end another --> it all depends on the machine &
the scheduler

```

**21/12/2022**

- Exercise:
  - Write the code of this using one statement:
    - One if-else
    - No loops



- Answer:

```
int main(){
```


```

if(fork()) //A
    fork() && (fork() || fork())
    //B      C      D
else fork() && fork()
    //E      F
}

```

• **Exercise:**

- Write the code of this with a condition
  - The parent doesn't create p2 before making sure that p1 created L1



- **Answer** (kind of):

```

int main(){
    int n=10,pid,s;
    for(int i=0;i<n;i++)
    {
        if(pid=fork()) //parent
            waitpid(pid,&s,NULL);
        else { //child
            fork(); //create child of child
            break;
        }
    }
}

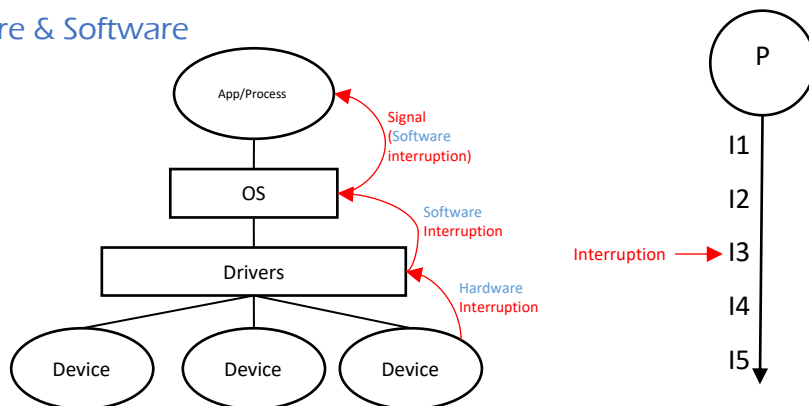
```

- **If I want all the children alive:** I can add an **infinite while loop** so that the children don't exit.

**4/1/2023**

- **Interruption**

- 2 types: Hardware & Software



- **Interruption vector**

- zone in **memory** related to the OS.
    - It is an **array of 2 fields**
      - **index**: identifies the interruption
      - **address of a handler**: a function that handles this interruption
  - **Context switching**: save the state of the process (everything inside the registers, variables..) → go to handler → perform the function → resume process
  - **Example**: executing I3 → interruption occurred → go to interruption vector to get address of handler → perform function → resume the process

Index	@ Of handler

- **Signal (Software Interruption)**

- **Definition**: a signal is a “software interruption” **delivered to a process**

- **Different sources**:

- **Internal** event (divide by zero, fault segmentation, Floating Point Error)
    - **External** event (Ctrl+C, Ctrl+Z)
    - **Explicit** request (I/O request)

- A signal reports the **occurrence of an exceptional event**

- Examples:

- Program errors: such as dividing by zero
      - User request to interrupt the process: Ctrl+C, Ctrl+Z
      - The termination of a child process (exit)
      - Expiration of a timer
      - A call to “kill” or “raise” command
        - **Kill** → process sends a **signal to another process**
        - **Raise** → process sends a **signal to itself**
      - Reading from an empty pipe with no writer

- **Each signal has an id or a macro**

- Library: **<signal.h>**
    - **List of signals command**: **Kill -l**

- id) name **(YOU DON'T HAVE TO KNOW ALL OF THEM)**

1) SIGHUP

2) SIGINT

3) SIGQUIT

4) SIGILL

5) SIGTRAP	6) SIGABRT	7) SIGBUS	8) SIGFPE
<b>9) SIGKILL</b>	<b>10) SIGUSR1</b>	11) SIGSEGV	<b>12) SIGUSR2</b>
13) SIGPIPE	<b>14) SIGALRM</b>	15) SIGTERM	16) SIGSTKFLT
<b>17) SIGCHLD</b>	<b>18) SIGCONT</b>	<b>19) SIGSTOP</b>	20) SIGTSTP
21) SIGTTIN	22) SIGTTOU	23) SIGURG	24) SIGXCPU
25) SIGXFSZ	26) SIGVTALRM	27) SIGPROF	28) SIGWINCH
29) SIGIO	30) SIGPWR	31) SIGSYS	34) SIGRTMIN
35) SIGRTMIN+1	36) SIGRTMIN+2	37) SIGRTMIN+3	38) SIGRTMIN+4
39) SIGRTMIN+5	40) SIGRTMIN+6	41) SIGRTMIN+7	42) SIGRTMIN+8
43) SIGRTMIN+9	44) SIGRTMIN+10	45) SIGRTMIN+11	46) SIGRTMIN+12
47) SIGRTMIN+13	48) SIGRTMIN+14	49) SIGRTMIN+15	50) SIGRTMAX-14
51) SIGRTMAX-13	52) SIGRTMAX-12	53) SIGRTMAX-11	54) SIGRTMAX-10
55) SIGRTMAX-9	56) SIGRTMAX-8	57) SIGRTMAX-7	58) SIGRTMAX-6
59) SIGRTMAX-5	60) SIGRTMAX-4	61) SIGRTMAX-3	62) SIGRTMAX-2
63) SIGRTMAX-1	64) SIGRTMAX		

#### ▪ Send a signal command

##### • Kill -sigid pid

- Sigid → signal id
- Pid → the process where I want to send the signal

##### • Example: kill -9 101 (sends signal SIGKILL to process with pid=101)

#### ◦ Receiving signals (3 options)

- Ignore the signal if the signal can be ignored (like SIGCHLD, SIGUSR1...)
- Do default action (action of the signal)
  - For most signals it is to kill the process, that's why we handle them
- Handle the signal if it can be handled
  - Similar to the handler in the OS, but here it is related to the process & we write how to handle it

#### ◦ 2 signals that I cannot ignore, and cannot handle:

- SIGKILL (9) → kill/end the process
- SIGSTOP (19) → pause the process not kill it (Ctrl+Z)
  - Resume the process: SIGCONT (18)

#### • Signal Handling:

- Handle the signal (instead of letting the default action take place)
- Handling function

##### ▪ Signal(sigID, handler)

- Handler is a **function I write** (it can be empty or do any other thing)
- the default action of the signal doesn't take place

```
signal(SIGFPE, myhandler);
void myhandler(int sig){ /*Anything Or Nothing*/ }
```

#### ◦ Signal function implementation

```
//how signal is written
typedef void (*sighandler_t) (int);
sighandler_t signal(int SIGNUM, sighandler_t action);
```

##### ▪ Action:

- SIG\_DFL (default) / SIG\_IGN (ignore) / Handler function
- returns the previous value of signal handler (the first return value is the default action)

#### • Handler:

- It takes an integer as argument. This integer is the code of the signal that triggered this function
- It is important if I want to handle different signals within the same handler function
- Example:

```
//Global
int counter = 0;
void myhandler(int sig){counter++;}

int main(){
    int i;
    signal(SIGCHLD, myhandler); //SIGCHLD --> signal sent by the OS when a child is created
    for(i=0;i<5;i++)
    {
        if(!fork()) exit(0);
        while(wait(NULL) != -1);
    }
    printf("Counter = %d\n",counter); //counter is 5
}
```

**9/1/2023**

- Sending signals:
  - To another process: `int kill(int pid, int sig)`
    - Pid: pid of the process to send the signal
    - Sig: signal id or name
  - To itself: `int raise(int sig)`
- Waiting signal:
  - `int pause()`
    - Blocking function
    - It stops the current process from work
    - Any other signal wakes the process
  - `int alarm(int sec):`
    - `sec` → waits `sec` times before sending a `SIGALRM` signal
    - doesn't block
    - if I want to block → add `pause()` while handling the `SIGALRM`\*\*\*\*\*
    - The first alarm to execute removes all the previous non-executed alarms

```
alarm(5);
alarm(3);
//since alarm(3) executes before alarm(5), alarm(5) will be canceled by alarm(3)
alarm(0); //cancel all alarms
```

- Signals `SIGUSR1` & `SIGUSR2` are only for communication. (no specific purpose)
- Exercise 1: Write a program that takes an input from the user. If the user doesn't input a value after 20sec, the program should print "You haven't entered a value yet" & ends the program.

```

void inputalarm(int signal)
{
    printf("\nYou haven't entered a value! (20s)");
    exit(); // stop waiting for input
}

void main()
{
    int x;
    signal(SIGALRM, inputalarm);
    printf("Enter a number: ");
    alarm(20);
    scanf("%d", &x);
    alarm(0); // cancel the alarm if the user enters a value
}

```

- **Exercise 2:** The SIGCHLD signal is sent to the parent when a child ends. Write the necessary code for a parent process that doesn't wait its child in blocking mode but the child won't become a zombie process (the parent has to wait when the child exits not before) (using signals)

```

void main()
{
    signal(SIGCHLD, zombiehandler);
    if (!fork())
    {
        sleep(3);
        exit(0);
    }
}

void zombiehandler(int signal){ wait(NULL);}

```

- **Exercise 3:** Parent that creates one child. The parent should print odd numbers between 1 and 100, while the child should print the even numbers between 1 & 100. The numbers should be printed in order.

```

void handler(int sig){}
int main(void){
    int child_pid;
    int parent_pid = getpid();
    signal(SIGUSR1,handler);

    if(child_pid = fork()){
        for(int number = 1; number <= 10; number += 2){
            printf("Parent: %d\n",number);
            kill(child_pid,SIGUSR1);
            pause();
        }
    }
    else{
        for(int number = 2; number <= 10; number += 2){
            pause(); // waiting for any type of signal
            printf("Child: %d\n",number);
            kill(parent_pid,SIGUSR1);
        }
    }
}
/*
Scenario:parent forks--> child in pause --> parent prints 1 & sends signal to child--> child awake, prints 2, sends
signal to child, then pauses --> parent wakes, prints 3, sends signal to child, pauses --> child wakes, prints 4,
send signal to parent, pause-->parent wakes.... */
Note: for a reason idk yet, the program doesn't work without \n in the printing

```

- **Exercise 4:** Write a program that displays on the screen alternating between tick & tock every 1 sec

```
int i=1;
void handler(int signal)
{
    printf("%s\n",i%2?"tick":"tock");
    i++;
}
void main(){
    signal(SIGALRM,handler);
    while(1){
        alarm(1);
        pause();
    }
}
```

- **Exercise 5:** Write a program that counts the signals it receives and displays the number of times each signal is received

```
int nsig[NSIG]; //NSIG --> number of signals present in the library <signal.h>
void handler(int signal){
    printf("Signal %d received %d times\n",signal,++nsig[signal]);
}
void main(){
    int s;
    for(s=1;s<NSIG;s++)
    {
        signal(s,handler);
        nsig[s] = 0;
    }

    while(1){
        pause();
    }
}
```

## **11/1/2023**

- **Exercise 1:** parent that creates N children, all paused. The parent sends for each child: SIGCONT, waits for 1s, send SIGSTOP until all children stop;

```
int i=0;
void Conthandler(int sig)
{
    printf("Child %i continued\n",i);
}
void main(){
    int pid[CHILDCOUNT];
    signal(SIGCONT,Conthandler);
    for(i=0;i<CHILDCOUNT;i++)
    {
        if(!(pid[i]=fork()))
        {
            pause();
            pause();
        }
    }

    if(i==CHILDCOUNT)
    {
        while(1){
            for(i=0;i<CHILDCOUNT;i++)
```

```

        {
            kill(pid[i], SIGCONT);
            sleep(1);
            kill(pid[i], SIGSTOP);
        }
    }
}

```

- **Exercise 2:** Ctrl+C is handled for 5 times, then it kills the program

```

int counter = 0;
void handler(int sig)
{
    counter++;
    printf("Ctrl+C handled\n");
}
void main(){
    signal(SIGINT, handler);
    while(counter < 5);
    signal(SIGINT, SIG_DFL);
    pause();
}

```

- **Exercise 3:** Parent creates a child, waits for the child to print its pid, then creates the next child.

```

void handler(int sig){}
void main(){
    signal(SIGUSR1, handler);
    for(int i=0; i<CHILDCOUNT; i++)
    {
        if(fork())
        {
            pause();
        }
        else
        {
            printf("Child pid = %d", getpid());
            kill(getppid(), SIGUSR1);
            break;
        }
    }
}

```

- **Exercise 4:** Parent creates children & waits for 5 seconds & then sends for the first child a signal & pauses. Child 1 wakes, child 1 sends signal to child 2, child 2 wakes, child 2 sends signal to child 3.... And so on until the last child wakes the parent. Use alarm. **Solution on page 32**

**16/1/2023**

## **Unix Shared Memory**

- Every process has its own address space.
- A **shared memory segment** is a zone that doesn't have pipe constraints. It is **shared by the processes that want to communicate with each other**.
- Created by one process (known as server) & attached it to its address space. This zone has an id.
- For the **client** (other process) to add it to its address space, it **calls the zone using its id**
  - this becomes the shared zone



- Problem: race conditions can occur
- Procedure for using shared memory:
  - Find a **key**. Unix uses this key for identifying shared memory segments
  - **shmget()**: to **allocate** a shared memory
  - **shmat()**: **attach** to shared memory space
  - **shmdt()**: **detach** shared memory space
  - **shmctl()**: **deallocate** shared memory
- Libraries to include:

```
#include <sys/types.h>
#include <sys/ipc.h>
#include <sys/shm.h>
```

- Keys:
  - of type **key\_t**
  - **global** entities
  - Do it yourself:

```
key_t somekey = 1234;
```

- Using **Ftok()**:

```
// ftok(char* path, int ID)
key_t somekey = ftok("./", 'a');
```

- Using **IPC\_PRIVATE**:
  - System provides it

- Allocate shared memory: **shmget()**

```
int shm_id = shmget(
    key_t KEY, //identity key
    int size,  //size of the shared memory
    int flag   //creation or use
);

key_t somekey = ftok("./", 'a');
shm_id = shmget(
    somekey, //key based on the current directory
    4*sizeof(int), //array of 4 integers
    IPC_CREAT | 0666 //creation | permits read & write
);
```

- Process that is **creating** the memory: flag = **IPC\_CREAT | 0666**
- Process that is **accessing** the memory: flag= **0666**
- Now, shared memory is allocated but not part of the address space

- Attaching shared memory: **shmat()**
  - For both server & client

- Returns a **pointer to the shared space**

```
int * shm_ptr = (int*) shmat(
    int shm_id; //return value of shmget()
    NULL,
    0
);
```

- You need to **cast the return value according to memory content**
  - In the above example: the memory contains an **array of 4 integers** → **shm\_ptr** is an **int \***
  - Use **shm\_ptr** to add, delete, & modify the shared data

- **Detach shared memory: shmdt()**

```
shmdt(shm_ptr);
```

- The shared space is **still there, but not attached to the memory**
- You **can reattach to it** using **shmat()**

- **Remove shared memory: shmctl()**

```
shmctl(
    shm_id,      //id returned from shmget()
    IPC_RMID,    //IPC remove ID
    NULL
);
```

- **If a parent wants to communicate with his child:** If parent created & attached the memory **before fork()**

- **Can use IPC\_PRIVATE** when creating shared memory
- **no need for shmget or shmat() in child**
- Because the child inherits this shared memory from the parent
- Example:

```
void Child(int data[])
{
    printf("%d %d %d %d\n",data[0],data[1],data[2],data[3]);
}

void main(int argc,char* argv[])
{
    int shm_id, *shm_ptr, status;
    pid_t pid;

    shm_id = shmget(IPC_PRIVATE,4*sizeof(int), IPC_CREAT | 0666);
    shm_ptr = (int*) shmat(shm_id,NULL,0);
    shm_ptr[0] = 1;
    shm_ptr[1] = 2;
    shm_ptr[2] = 3;
    shm_ptr[3] = 4;

    if((pid=fork()) == 0)
    {
        Child(shm_ptr);
        exit(0);
    }

    wait(&status);
    shmdt((void*)shm_ptr);
    shmctl(shm_id,IPC_RMID,NULL);
    exit(0);
}
```

- **Communicate between 2 separate processes:**
  - Need a key, **shmget()**, **shmat()**, & **shmdt()** for both processes

- Need only 1 shmctl()
- Key should use ftok() with the same parameters
- Server must run first to prepare the shared memory
- Example:
  - This version uses busy waiting

```
#define NOT_READY -1
#define FILLED 0
#define TAKEN 1

struct memory{
    int status;
    int data[4];
};

//Server
void main(int argc, char*argv[])
{
    key_t shm_key;
    int shm_id;
    struct memory * shm_ptr;

    //prepare shared memory;
    shm_key = ftok("./", 'x');
    shm_id = shmget(shm_key, sizeof(struct memory), IPC_CREAT | 0666);
    shm_ptr = (struct memory*) shmat(shm_id, NULL, 0);

    //Fill memory
    shm_ptr->status = NOT_READY;
    for(int i=0; i<4; i++) shm_ptr->data[i] = i;
    shm_ptr->status = FILLED;

    //Wait for child to be done
    while(shm_ptr->status != TAKEN) sleep(1);

    //detach & remove
    shmdt((void*)shm_ptr);
    shmctl(shm_id, IPC_RMID, NULL);
    exit(0);
}

//Client
void main(int argc, char*argv[])
{
    key_t shm_key;
    int shm_id;
    struct memory * shm_ptr;

    //Prepare shared memory
    shm_key = ftok("./", 'x');
    shm_id = shmget(shm_key, sizeof(struct memory), IPC_CREAT | 0666);
    shm_ptr = (struct memory*) shmat(shm_id, NULL, 0);

    //Wait for parent to fill
    while(shm_ptr->status != FILLED);

    //print
    for(int i=0; i<4; i++) printf("%d ", shm_ptr->data[i]);
    shm_ptr->status = TAKEN;

    //detach
    shmdt((void*)shm_ptr);
    exit(0);
}
```

- **Note:** If you didn't remove your shared memory segments using shmctl(), they will be in the system forever → degrade system performance
  - lpcs command: check if you have left shared memory segments
  - lpcrm: remove shared memory segments

- **Exercise:** write a program where 2 processes communicate using shared memory. P1 fills memory with n integers. P2 calculates the average of these values. P1 prints this value.

```
void main(){
    int i,pid, seg;
    int N = 3;

    //creating & attaching the shared zone
    key_t Key1 = ftok("./",'a');
    seg = shmget(Key1, sizeof(int)*(N+1),IPC_CREAT | 0666);
    int * data = (int*) shmat(seg,NULL,0);

    //child calculates average
    if( (pid = fork()) == 0)
    {
        while(data[N-1] == 0) sleep(3); //child needs to wait for the parent to fill the array
        printf("Child: calculating average!\n");
        data[N] = 0;
        for(int i=0;i<N;i++) data[N] += data[i];
        data[N] /= N;
        exit(1);
    }
    else{
        srand(getpid()); //for randomizing the seed
        printf("Parent: filling array!\n");
        for(int i=0;i<N;i++)
        {
            sleep(rand()%3);
            data[i] = rand()%1000;
        }

        for(int i=0;i<N;i++) printf("%i ",data[i]);
        printf("\n");
        wait(NULL);

        printf("Parent: average is %d\n",data[N]);
        shmdt(0);
        shmctl(seg,IPC_RMID,0);
        printf("END\n");
    }
}
```

**18/1/2023**

- **Previous exercise 4:** Parent creates children & waits for 5 seconds & then sends for the first child a signal & pauses. Child 1 wakes, child 1 sends signal to child 2, child 2 wakes, child 2 sends signal to child 3.... And so on until the last child wakes the parent. Use alarm.
  - Answer:

```
//Exercise 4
int pid,fd[2],ctr,prid_main;
void handler(int signal){
    ctr++;
    if(signal==SIGALRM)
    {
        read(fd[0],&pid,sizeof(int));
        printf("Parent woke & sent signal to child with pid: %d\n",pid);
        kill(pid,SIGUSR1);
        pause();
    }
    else if(signal == SIGUSR2)
    {
        printf("Parent %d exited\n",getpid());
        exit(0);
    }
    else
    {

```

```

        if(read(fd[0],&pid,sizeof(int)))
        {
            printf("Child %d sent signal to child %d\n",getpid(),pid);
            kill(pid,SIGUSR1);
        }
        else //last child
        {
            printf("Last child %d sent signal to parent %d\n",getpid(),getppid());
            kill(getppid(),SIGUSR2);
        }
        exit(0);
    }
}

void main(){
    int N=5;
    pipe(fd);
    signal(SIGUSR1,handler);
    signal(SIGUSR2,handler);
    signal(SIGALRM,handler);

    for(int i=1;i<N;i++)
    {
        if(!(pid=fork()))
        {
            close(fd[1]);
            pause();
        }
        else
        {
            printf("Process %d created child with pid: %d\n",getpid(),pid);
            write(fd[1],&pid,sizeof(int));
        }
    }

    close(fd[1]);
    alarm(5);
    pause();
}

```

## Memory Management

- Processes are loaded in memory
- **Logical address vs physical address:**
  - **Logical: virtual address** that a process generates for accessing memory. They are usually expressed in terms of the size of the word in the computer's memory. Operating systems use logical addresses to map it to a physical address.
  - **Physical address:** a **memory** address that refers to a specific location in the physical memory.
- Questions we ask:
  - Where are processes loaded?
  - Do we load all the programs into memory or part of programs?
  - What if the size of the program > memory size?
  - How to store several processes whose total size > memory size?
  - How to manage free spaces & used spaces in memory?
  - How to protect the address space of the process in memory?
  - How to share data between processes in memory?

- Definition: the act of managing computer memory
  - Providing ways to **allocate portions of memory to programs at their request**, and freeing it when no longer needed.
  - **It is a linear array of bytes where each byte is named by a unique memory address**



- Recall: processes are defined by an address space (the addresses that are accessible by the process (code, data, heap, & stack))
  - We cannot know where a program will be loaded ahead of time


- **Address Binding:**

- Definition: fixing physical address to logical address of a process' address space.
- Types:
  - Compile time binding: if program location is fixed & known ahead of time
    - We identify the address during compilation
  - What we'll focus {
    - Load time binding: program location in memory is unknown until run-time and location is fixed
    - Execution time binding: process can be moved in memory during execution
      - Requires hardware support!

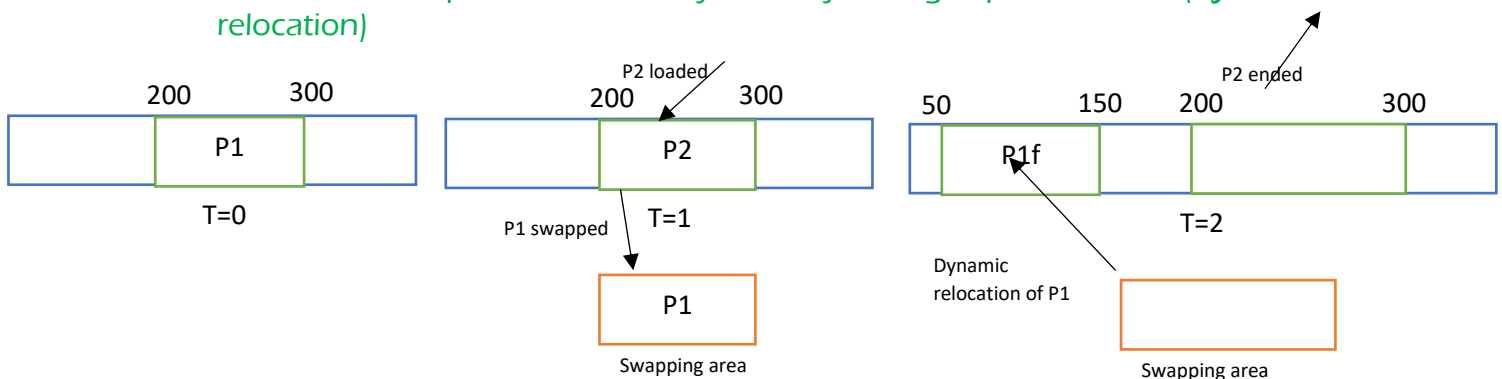
- **History of Memory Management Schemes:**

- **No memory abstraction:** address binding
  - Only **one program can be run**
  - **Physical memory is allocated from 0 to some max**
  - Running several programs at the same time **requires swapping**
    - Swapping: there exists a swap area in the hard disk where processes are swapped to be executed.
- **Memory Abstraction:**
  - **Each process has its own address space**
  - During compilation, we give the **process logical addresses from 0 to max.**
  - **Physical address = base + logical address (relocation)**
    - Base register stores the base value
    - Limit register stores the limit/max value

- The physical address ( $\text{base} + \text{logical address}$ ) should be  $< \text{limit}$  (protection)
- In this case: physical address = 200 + logical address




- Logical addresses are from 0 to 100
- Base register: 200
- Limit register: 300
- Each process has its base register & its limit register.
- The location of a process in memory can vary during implementation (dynamic relocation)



### Dynamic Relocation

- Memory Management Unit (MMU) is the hardware device that converts logical to physical addresses based on the schema



- Multiprogramming:



- Many processing being executed in memory
- When a program is running, the **entire program must be in memory**
- Memory is divided into partitions
- Each program is put into a single partition

- Swapping:

- Occurs between processes to allow multiple programs to be executed at once
- Processes are swapped in & out of memory

- Fragmentation:

- Occurs when swapping is made
- Free spaces (holes) in memory distributed in the memory. Each hole is **not sufficient to hold an entire process** → lost space & degradation of performance.
- 2 types of fragmentation:
  - **External fragmentation:** when the size of the partitions are variable according to the processes' size
  - **Internal fragmentation:** when I divide the continuous layout into partitions of fixed size (explained later)
- Example:




- In here we have external fragmentation

- How to deal with external fragmentation:


- Compaction

- group processes beside each other, and free spaces beside each other
- Disadvantage: takes too much time






- Partition sizes:
  - Programs grow during execution (more room for stack, heap allocation,...)
  - Problem:
    - If partition too small → programs must be moved → copying overhead & a lot of swapping
  - Temporary Solution: allocate a bit of extra space for program growth 😞
- Management Data Structures:
  - How to identify used & unused spaces?
  - A chunk of memory is either used or unused (free)
  - Bit Maps:
    - 1<sup>st</sup> data structure to keep track of used & unused memory
    - It is a long string of bits, where each bit represents a chunk of memory.
      - Bit is 1 → used space
      - Bit is 0 → unused space



- Linked List:
  - 2<sup>nd</sup> data structure to keep track of used & unused memory
  - We have a struct consisting of:
    - Bit indicating the type of space (P → process, H → hole (free))
    - Starting address
    - Length
    - Pointer to next element
  - A linked lists of elements is saved that gives a map of the memory spaces.




- Algorithms used to load processes:

- If I have multiple free spaces, where do I store the process in memory?
- **First Fit Algorithm**: start from the **beginning of the memory** and load the process at the **first free zone** that is equal or bigger than the size of the process
- **Best Fit Algorithm**: Search the **entire memory from beginning to end**, and take the **smallest hole that is adequate**
- **Worst Fit Algorithm**: Search the entire memory from beginning to end, and take the **largest hole that is adequate**
- **Next Fit Algorithm**: Like the **First Fit Algorithm**, but it searches from the last placed process (not from the beginning of the memory)

- Memory Layout scheme:

- Memory is made up of a set of consecutive bytes (addresses)
- **Continuous layout**:
  - **Layout 1**: the process is loaded entirely into the memory as a consecutive bunch.
    - **Variable sized partitions** according to the size of the process (like previously explained)
    - If no free space for a new process → **Swapping**
    - Problems:
      - overhead in writing to and reading from disk
      - External fragmentation (solved by compaction)
      - How to execute processes whose size is greater than memory size?
      - No sharing processes is possible
  - **Layout 2**: divide the memory into fixed partitions with different sizes
    - Each partition is associated with a queue of processes.
    - Each process is put in the minimum partition it can fit in.
    - Example:




- processes  $\leq 50$  KB → load it into A
- processes  $> 50$ KB & processes  $\leq 100$ KB → load into B
- processes  $> 100$ KB & processes  $\leq 200$ KB → load into C
- Processes  $> 200$ KB & processes  $\leq 300$ KB → load into D

- **Advantage**: solves the problem of external fragmentation

- Suppose this scenario for the above example:


- P1: 30KB & P2: 40KB arrive
- P1 is loaded in partition A & P2 is enqueued in the same partition



A queue: P2


→ Free space of 20KB

- P1 swapped out (overhead) and P2 is loaded in the memory



→ Free space 10KB in A

- Suppose P3: 20KB arrived
- P2 is swapped and P3 is loaded in the memory




→ Free space 30KB in A

- Disadvantage: Internal fragmentation (free space inside fixed size partitions)


#### ○ Non-continuous layout schemes: (paging)

- Based on virtual memory
- Associate to each process an address space (virtual) equal to the max addressing size
- Example: if I have an architecture: 32 bits → it can hold a max of  $2^{32}$  bits (4GB)



#### 1. Divide the address space into equal and fixed chunks called pages (paging)

- The size of the page is a pre-determined parameter configured during the installation of the system (usually 4KB)
- It should be a power of 2 ( $4KB = 2^{12}$ )
- Suppose a process with size 16KB
- Process → set of consecutive pages from 0 to max



- Number of pages =  $\frac{\text{size(process)}}{\text{size(page)}} = \frac{16\text{ KB}}{4\text{ KB}} = \frac{2^4 \times 2^{10}}{2^2 \times 2^{10}} = \frac{2^{14}}{2^{12}} = 2^2 = 4\text{ pages}$


- **Logical address** now consists of 2 parts (**page number**, **offset**)
  - Offset: placement inside the page
  - How to find the values
    - **Page number:**  $\frac{\text{address}}{\text{size}(\text{page})}$
    - **Offset:**  $\text{address} \% \text{size}(\text{page})$
  - Page 0 contains addresses from 0 to 4KB → from 0 to  $2^{12}$  B → 0 to 4095 B
    - 1<sup>st</sup> address(0, 0), 2<sup>nd</sup> address (0,1) ... until last address (0,4095)
  - Page 2 contains the next 4KB → from 4096 to 8191....
    - 1<sup>st</sup> address is (1,0), 2<sup>nd</sup> is (1,1) ... until last address (1,4095)

▪ **Example:** CPU wants to access to address 200 B

- Which page? Which offset?
- Page:  $\frac{200}{\text{size}(\text{page})} = \frac{200}{4096} = 0$
- Offset:  $200 \% \text{size}(\text{page}) = 200 \% 4096 = 200$
- Logical address → (0,200)

2. Divide the physical memory into fixed and equal chunks called **frames**

- $\text{size}(\text{frame}) = \text{size}(\text{page})$




- **Physical address** now consists of (**frame number**, **offset**)
- **Physical address = (frame position \* size(offset)) + offset**

3. The process isn't loaded entirely into the memory.

- It is loaded **page by page** as requested.
- The pages of the process in memory should not be consecutive

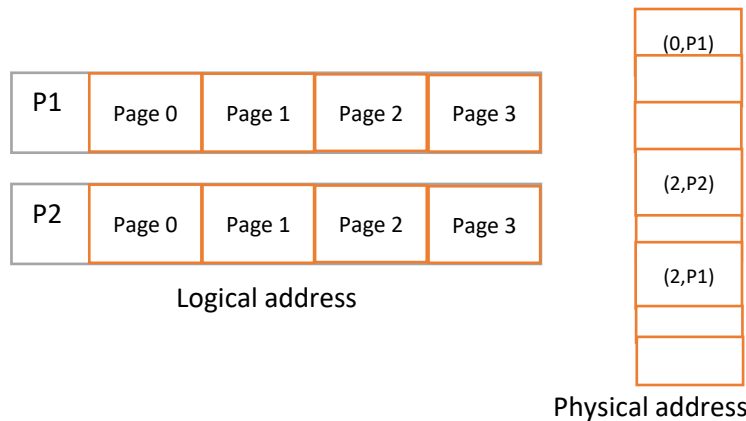
4. How to load the process in memory?

- The operating system **decomposes the address space of the process as set of pages**.
- These pages are stored consecutively in the swap area.



- **During execution:**

- CPU makes access to address 255 → (0,255)
- Load the page that contains that @ 255 into memory (page 0)
- The memory manager stores the loaded pages into one of the free frames



iv. The **logical address is not the same as the physical address**

- However, offset in physical @ = offset in logical @

- **Example:** address 8112 in P2

- a. Logical address is  $(8112/4KB, 8112 \% 4KB) = (2,0)$

- b. In physical address, page 2 was placed in frame 4 (index 3)

- c. Physical address =  $(4 * 4096) + 0 = 16384$

- So, we have a translation from logical address (pg #, offset) to physical address (frame #, offset)

- i. Offsets are the same

- ii. Page numbers are different

- **Example 1:** if I have an architecture of 20 bits addressing with 4KB page size

- Questions:

- How much addresses exist?
    - How many bits are needed to encode the page number?
    - How many pages can we have?

- Answers:

- The address is 20 bits. So, we can have  $2^{20}$  different addresses
    - Each page size is  $4KB = 2^{12}$  B. So each page contains  $2^{12}$  addresses → I need 12 bits to represent the offset. Therefore, I need  $20 - 12 = 8$  bits to represent the page number in the address
    - Pages are encoded in 8 bits → I can have  $2^8 = 256$  pages

- **Example 2:** if I have an architecture of 32 bits addressing with 4KB page size

- Questions:

- How much addresses exist?
    - How many bits are needed to encode the page number?
    - How many pages can we have?

- Answers:

- I can have a total of  $2^{32}$  different addresses
    - $\text{Size}(\text{page}) = 2^{12} \rightarrow 2^{32}/2^{12} = 2^{20}$  total pages → I need 20 bits to represent the pages.
    - $2^{20}$  pages

- **Example 3:** access address 32780 in a 32 bit address architecture and page size of 4KB

○ I have to write the address in the form of (page#, offset)

○ **Way 1:**

▪ Convert the address from decimal to binary

• 0000 0000 0000 0000 1000 0000 0000 1100

• The offset is 12 bits (calculated in example 2) → in decimal: 12

• The page # is the rest (20) → in decimal: 8

○ **Way 2:**

▪ Page #:  $\text{Address} / \text{size}(\text{page}) = 32780 / 4096 = 8$

▪ Offset:  $\text{Address} \% \text{size}(\text{page}) = 32780 \% 4096 = 12$

○ → (8,12)

• **Conversion from logical address to physical address**

○ To know where and if a page is loaded in memory, the MMU uses a page table to show the mapping of the virtual address to the physical address.

○ Each process has a page table associated with it.

○ The memory address of the page register is stored in a special register

○ The page table has entries = the maximum number of pages

○ Each Page Table Entry (PTE) is 32 bits

○ Attributes of the page table:

▪ **Presence/absence bit**

• 1 → page is loaded in physical memory & can be used (page hit)

• 0 → page is not loaded in physical memory (page fault)

▪ **Protection bit**

• The permissions

• If on 1 bit:

○ 0 → read & write

○ 1 → read only

• If on 3 bits:

○ rwx (read, write execute) → a bit for each

▪ **Modified**

• Dirty → it must be written to disk

• Clean → it can be rejected

▪ **Referenced**

• Useful in choosing the victim for removing when I have a page fault & the memory is full

• It is the number of times the loaded page was accessed.

• We choose to remove the page with minimum references.


▪ **Frame number**

• The number of frame in physical memory the page is loaded in


Presence	Protection	Modified	Referenced	Frame #

○ **Example:** Consider an address space  $64\text{KB} = 2^{16}$  &  $\text{size}(\text{page}) = 4\text{KB} = 2^{12}$  & physical memory of  $32\text{KB}$

- Address space =  $2^{16} \rightarrow$  I am working in 16 bit architecture
- $\text{Size}(\text{page}) = 2^{12} \rightarrow$  I need 12 bits for the offset
- Number of bits for the page number:  $16 - 12 = 4$
- Number of pages:  $2^4 = 16$
- $\text{Size}(\text{memory}) = 32\text{KB} \rightarrow$  I have  $32\text{KB} / 4\text{KB} = 8$  frames



- Suppose the CPU makes reference to the address  $200\text{ B} \rightarrow (0, 200)$ 
  - The MMU checks that the presence bit of page 0 is 0  $\rightarrow$  page fault
  - The MMU loads the page 0 into memory in any free space
  - Suppose it loaded it in frame number 7
  - $(0, 200) \rightarrow (7, 200)$
  - The physical address will be:  $(7 * 4096) + 200$




- **Question:** in 64 bit architecture &  $\text{size}(\text{page}) = 4\text{KB}$ 
  - How many entries are contained in the page table?
  - Calculate the size of the page table?
- **Answers:**
  - Entries in page table = number of pages.


- $\text{Size}(\text{page}) = 2^{12} \rightarrow 12 \text{ bits for the offset} \rightarrow 64 - 12 = 52 \text{ bits for the page number} \rightarrow 2^{52} \text{ pages}$
- $\text{Number of PTE} = 2^{52}$ 
  - Each entry is 32 bits  $\rightarrow \text{size}(\text{page table}) = 2^{52} * 32 - 2^{53} * 25 = 2^{57}$
- **Problems we solved with paging:**
  - Storing pages with sizes  $> \text{size}(\text{memory})$
  - Store several processes with size  $> \text{size}(\text{memory})$
  - Protect address space of a process in memory
  - Share data between processes in memory (sharing pages)
  - Solved external fragmentation
- **Small problem arises:**
  - Internal fragmentation
  - Example: a process of size 9KB &  $\text{size}(\text{page}) = 4\text{KB}$ 
    - The process will fill 2 pages & use 1KB from the 3<sup>rd</sup> page
    - The 3KB not used in the 3<sup>rd</sup> page is what we call internal fragmentation
- **Multi-level page table:**
  - Why do I need a multi-level page table?
    - The page table is stored in the memory  $\rightarrow$  it needs frames to store the whole table. However, **processes usually don't use all of the PTEs** inside the page table (they only access some). So it would be a **waste of space to import all the page table into memory**
  - Example:
    - On 32bit architecture & 4KB page size
      - Max size of process:  $2^{32} = 4\text{GB}$
      - $2^{20}$  different pages
      - $2^{20}$  PTE
      - Page table size:  $2^{20} * 32 = 2^{25} = 4\text{MB}$
      - Storing the whole page table in memory:
        - I need 1024 frames (4MB/4KB)
      - Suppose we have a process 14MB:
        - Number of pages this process use:  $14\text{MB}/4\text{KB} = 14 * 2^8 \ll 2^{20}$
      - So, there is a lot of empty pages stored inside the memory
    - So, we use multi-level paging, aka paging of the page table, where we only store the pages we need in the memory.
  - Now, to **access the physical memory, we need to pass through different page tables.**
  - The initial page table that we have is paged
    - We **divide it into pages and create a new page table for the page table**
    - **Example:**



- Previously:



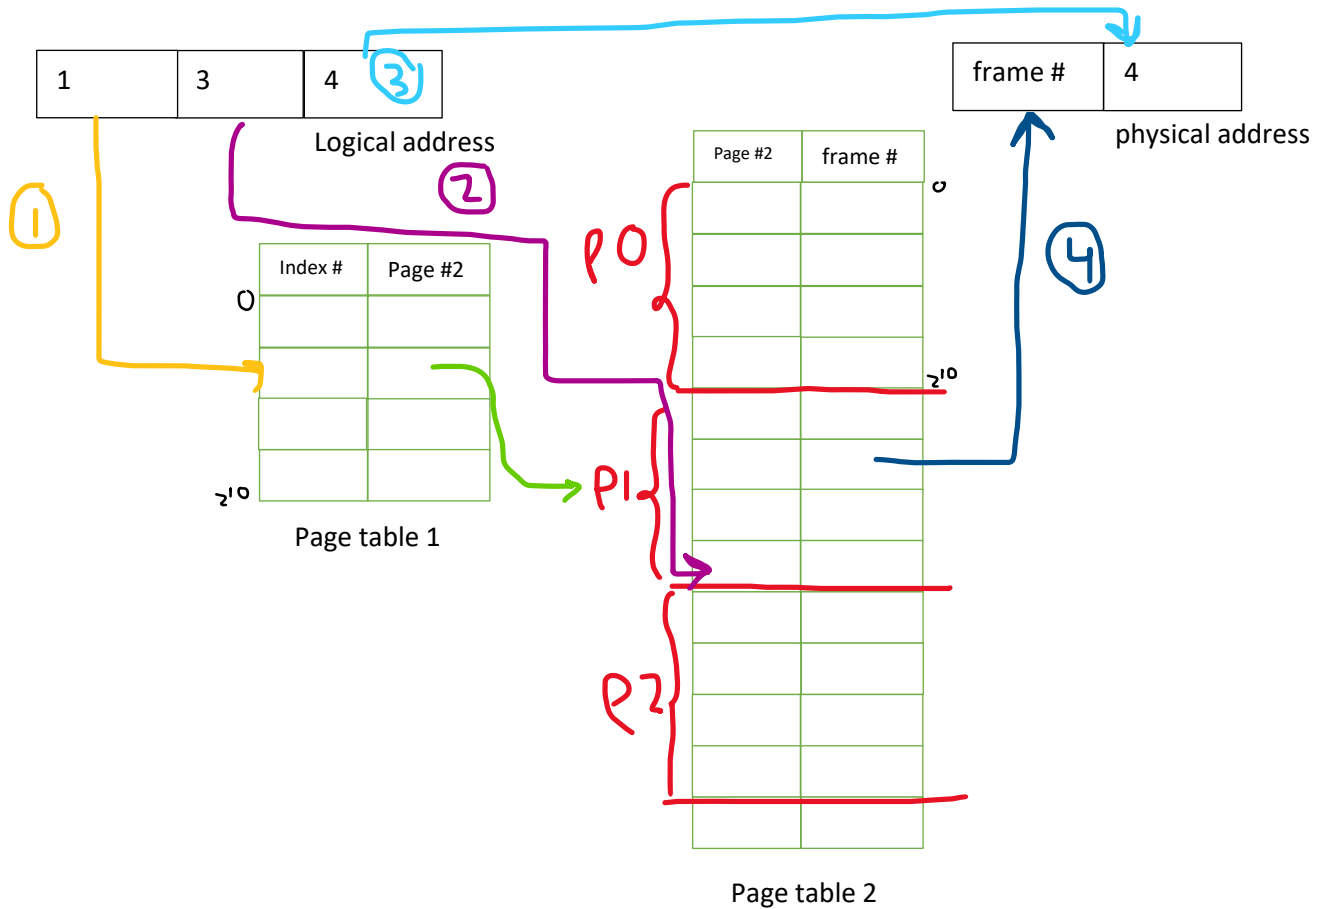
- Now:



- Accessing the physical memory is divided into 2 parts
  - **First level (outer table)**
    - To indicate which chunk in the page table is the address in
  - **Second level (inner table)**
    - To indicate which frame the page is stored in memory
- **Physical address = (index # \* size(chunk)) + (page # \* size(page)) + offset**
- **Example:**
  - In a 32bit address & 4KB page size:
    - In a single-level page table: I would need 20 bits for the page table, & 12 for the offset

- In a multi-level page table: I would need 10 bits for the index number, 10 bits for the page number, and 12 bits for the offset
- How did I know that I need 10 bits for the index:
  - PTE is 32 bits = 4B
  - Page size is: 4KB
  - I need to divide the page size by the PTE:  $4\text{KB}/4\text{B} = 2^{10}$  entries → The page table is divided into  $2^{10}$  pages/chunks
  - I need 10 bits to represent each chunk in the page table
- Example: CPU to move to register address 4,206,596 B (32bit architecture & 4KB page size)
  - We know that we have 10 bits for index number & 10 bits for page number & 12 for offset
  - Way 1:
    - Change decimal to binary → 0000 0000 0100 0000 0011 0000 0000 0100
    - Index #: 0000 0000 01 → 1
    - Page #: 00 0000 0011 → 3
    - Offset: 0000 0000 0100 → 4
  - Way 2:
    - 10 bits for index → we have  $2^{10}$  chunks/indexes
    - Each chunk has  $2^{10}$  pages where each page size is  $2^{12}$
    - Each chunk size is  $2^{10} * 2^{12} = 2^{22}$  B = 4MB
    - Each chunk has  $2^{22}$  addresses
      - First chunk: 0 → 4194303
      - Second chunk: 4194304 → 8388607 ....
    - So, the address 4206596 is in chunk:  $4206596/2^{22} = 1^{\text{st}}$  chunk
    - What page is it in the page table?
      - $4206596 - (1 * 2^{22}) = 12292$  address in  $2^{\text{nd}}$  page table where each page size is  $2^{12}$
      - $12292/2^{12} = 3^{\text{rd}}$  page
    - Which offset?  $12292 - (3 * 2^{12}) = 4$

$$4206596 = (1 * 2^{22}) + (3 * 2^{12}) + 4$$



#### Problems of Multi-level:

- Slows down the performance/speed of CPU by a lot

#### Solution: Inverted Page Table

#### Consider this scenario:


- System architecture 64 bits
- Page size 4KB
- Memory size 512 MB
- Questions:

- How much space would a single-level page table take?
- How much space would a multi-level page table take?

#### Answers:

- Page size 4KB =  $2^{12}$  B  $\rightarrow$  I need 12 bits for offset  $\rightarrow$  52 bits (64-12) for page number. So, I have  $2^{52}$  PTEs where each page is 32 bits = 4B. Size of single level page table is:  $2^{52} * 2^2 = 2^{54}$  B

- $\text{Size}(\text{page}) = 4\text{KB} = 2^{12} \text{ B}$  &  $\text{size}(\text{PTE}) = 32 \text{ bits} = 4\text{B}$ . I need  $2^{12}/2^2 = 2^{10}$  entries for each multi-level table. If I have 12 bits for the offset, the other 52 bits should be divided into 10bit parts.




- $\rightarrow$  6 levels on indexing  $\rightarrow$  each address translation needs 6 memory references  
 $\rightarrow$  performance problem!!!

### • Inverted Page table

- Make a **single and global page table related to the physical memory** and not related to a process.
- Consider the example of the above scenario
  - $\text{Size}(\text{memory}) = 512 \text{ MB} = 2^{29} \text{ B}$
  - $\text{Size}(\text{page}) = \text{size}(\text{frame}) = 4\text{KB} = 2^{12} \text{ B}$
  - Number of frames I need is:  $2^{29}/2^{12} = 2^{17}$  frames
  - Create a table with  $2^{17}$  entries
  - Each entry is related to a physical frame in memory
  - Suppose the entry is 16B  $\rightarrow$  the size of the whole page table is  $2^{17} * 2^4 = 2\text{MB}$

### ◦ The inverted table includes:

- The **presence/absence field, modified field, protection field** (same as normal page table)
- **Pid & page number of the process loaded in this frame**




### ◦ How is the logical address converted?

- For each logical address, decompose it as usual (**page #, offset**) & we have the **pid of the process**
- The **MMU searches in the inverted page table** in either 2 ways
  - **Linear:** It searches the page table **entry by entry**
    - If the entry is found (page hit): I get the frame number

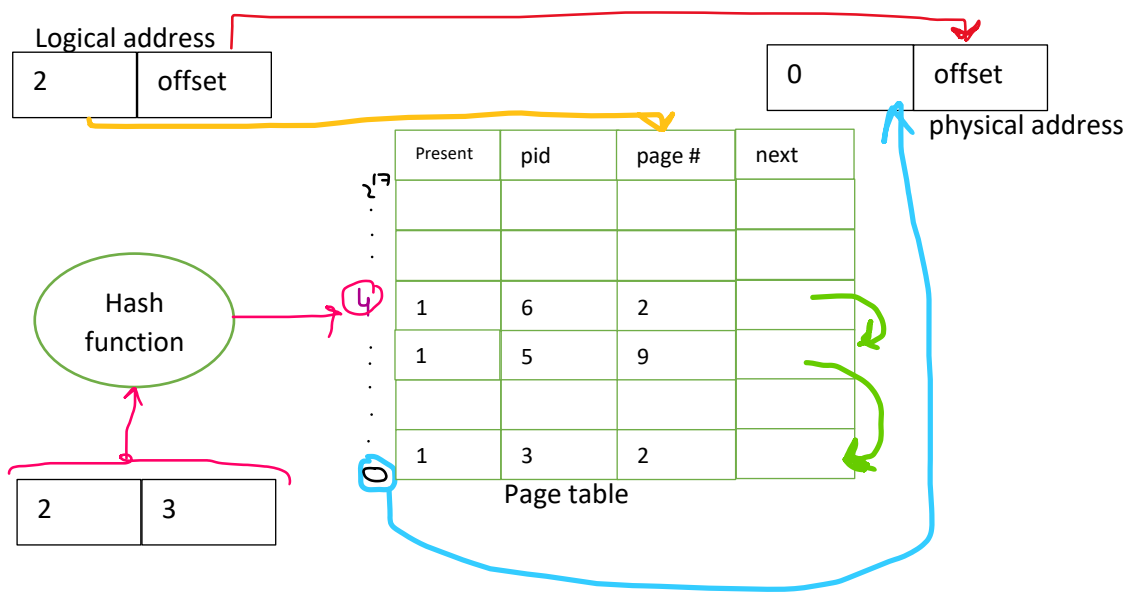
- If the entry is not found (page fault): load page from disk and update the corresponding entry in the inverted page table

- **Hashing:**

- We use a **hash function that hashes a key into a value**
- This value will give the index of the frame in the inverted page table
- For the same key, the hash function outputs the same value.



- If 2+ keys have the same value, collision occurs. We can use a linked list to keep track of the keys with the same value.
- Example with collision:




## • Segmentation

- The process is made up of data, code, heap, & stack.
- Paging uses one continuous sequence of all virtual addresses from to the maximum needed for the process
- Segmentation is an alternate schema that uses **multiple separate address spaces for various segments of a program. (process is divided into segments)**
- A **segment** is a **logical entity** in which the program is aware of (function, array, stack...)
- Each **segment** has a **variable length** that may change during execution
- The **logical address is (segment #, offset)**
  - Segment # to specify which segment the address is located in
- It is **similar to dynamic partitioning** (because of variable sizes). However, the difference is that **segments can be not contiguous**.
- It eliminates internal fragmentation, but **suffers from external fragmentation**
- There is no relationship between logical & physical address
- **A segment table is needed. It includes:**
  - **Segment number**
  - **Base:** The base/starting address of the segment
  - **Length:** The length of the segment



- **Example:** Suppose I want to find c with logical address 0x02
  - Which segment? Segment number 1 has a length of 6 → from 0x00 to 0x05 → c is in segment 1
  - Since it is in segment 1, the offset is 2
  - Logical address → (1, 2)
  - The segment table indicates that segment 1 starts at physical address 0x02. But the offset of c is 2
  - Physical address of c is  $0x02 + 2 = 0x04$
- Segments **allow sharing between processes** (like sharing libraries)
  - The library can be put in a segment and shared by multiple processes
  - Avoiding the need to put the library in each process's address space
- The **user can put protection on each property/segment**

- A function segment can be set to execute but not read or write.



- **Segmentation with paging**
  - The process is divided into segments
  - Each segment is divided into pages
  - Logical address  $\rightarrow$  (segment #, page #, offset)
  - Physical address  $\rightarrow$  (frame # \* size(page)) + offset
- **Exercise:** memory with segmentation & paging
  - Page size: 4KB
  - Physical memory: 64KB
  - Process P is composed of 3 segments: S1 (16KB), S2 (4KB), S3 (8KB)
  - At instance t
    - Page 1 Segment 1 loaded into physical memory 2
    - Page 2 Segment 1 loaded into physical memory 0
    - Page 0 Segment 2 loaded into physical memory 9
    - Page 0 segment 3 loaded into physical memory 12
  - Questions: For decimal address 8212, Indicate the
    1. segment number
    2. page number in the segment
    3. offset
    4. frame number
    5. physical address
  - Answers:
    1. # of frames in memory:  $64\text{KB}/4\text{KB} = 16$ 
      - Segment 1 is 16KB: from address 0  $\rightarrow 2^{14}$  (16384)  $\rightarrow$  address 8212 is in **segment 1**
    2. # of pages in segment 1:  $16\text{KB}/4\text{KB} = 4$  each with size 4KB

- First page from 0 → 4095
  - Second page from 4096 → 8191
  - Third page from 8192 → 12287....
  - So address 8212 is in the 3<sup>rd</sup> page → **page # = 2**
3. Page # 2 starts with address 8192 → offset: 8212 – 8192 = 20
4. Page 2 Segment 1 → **frame = 0** (from given)
5. **Physical address = (frame # \* size(frame)) + offset**
- Physical address = (0 \* 4096) + 30 = 20

#### • Page Replacement Algorithms

- If the **memory is full & I need to load a new page** into the memory, I need to evict a page in the memory to replace it with the new one.
- Steps:
  - Know which page I need to remove
  - If it has been modified, write it back to disk
  - Load the page into memory
- **Which page do I evict?** I have to leave in memory the most important pages → the most requested pages. This is because I need to avoid excessive loading from disk.
- There are 2 strategies to find which page:
  - Local: to evict the pages that are related to this process
  - Global: to evict pages with respect to the whole memory (Our course of study)
- There are **multiple algorithms that help us choose which page to choose**.
- **Terminology:**
  - **Reference String:** the memory reference sequence generated by a program
  - **Paging:** moving pages to/from disk
  - **Optimal:** best (theoretical) strategy
  - **Eviction:** throwing something out
- Algorithms:
  - **Random Algorithm:** choose a page **randomly**
  - **Optimal Algorithm:**
    - **Theoretical** (never used), It is only for comparison purposes
    - Assumes that we **know the future reference string for a program** → It would **choose the page with the last reference in the future**

Page Refs	3 Page Frames		
	Fault?	Page Contents	
A	yes	A	
B	yes	B	A
C	yes	C	B A
D	yes	D	B A
A	no	D	B A
B	no	D	B A
E	yes	E	B A
A	no	E	B A
B	no	E	B A
C	yes	C	E B
D	yes	D	C E
E	no	D	C E

A & B will be called the most recent

A & B will be called the most recent

A will not be called

B will not be called

7 faults



▪ FIFO:

- The first page in is the page to go

Page Refs	3 Page Frames			
	Fault?	Page Contents		
A	yes	A		
B	yes	B	A	
C	yes	C	B	A
D	yes	D	C	B
A	yes	A	D	C
B	yes	B	A	D
E	yes	E	B	A
A	no	E	B	A
B	no	E	B	A
C	yes	C	E	B
D	yes	D	C	E
E	no	D	C	E

9 faults

- **Anomaly:** As the number of page frames increases, the number of faults increases

Page Refs	4 Page Frames				
	Fault?	Page Contents			
A	yes	A			
B	yes	B	A		
C	yes	C	B	A	
D	yes	D	C	B	A
A	no	D	C	B	A
B	no	D	C	B	A
E	yes	E	D	C	B
A	yes	A	E	D	C
B	yes	B	A	E	D
C	yes	C	B	A	E
D	yes	D	C	B	A
E	yes	E	D	C	B

10 faults

▪ Least Recently Used:

- Removes the page that has not been referenced for the longest time
- Keep track of the  $t_{\text{last accessed}}$
- Reference string: {A, B, C, D, A, B, E, A, B, C, D, E}

	Fault	F1	F2	F3
A $t_{\text{last accessed: 0}}$	Yes	A		
B $t_{\text{last accessed: 1}}$	Yes	B	A	
C $t_{\text{last accessed: 2}}$	Yes	C	B	A
D $t_{\text{last accessed: 3}}$	Yes	D	C	B
A $t_{\text{last accessed: 4}}$	Yes	A	D	C
B $t_{\text{last accessed: 5}}$	Yes	B	A	D
E $t_{\text{last accessed: 6}}$	Yes	E	B	A
A $t_{\text{last accessed: 7}}$	No	E	B	A
B $t_{\text{last accessed: 8}}$	No	E	B	A
C $t_{\text{last accessed: 9}}$	Yes	C	B	A
D $t_{\text{last accessed: 10}}$	Yes	D	C	B
E $t_{\text{last accessed: 11}}$	Yes	E	D	C

10 faults

▪ According to the number of references & dirty bit

- **Reference bit:** increases every time the page is read or written
- **Dirty bit:** set when page is written to

- Cases:

Reference	Dirty Bit
0	0
0	1
1	0
1	1

- Remove the page that is least referenced (R=0) & not written to (d = 0) (else the one with R=0 & D=1, else R=1 & D = 0...)

- Second Chance Algorithm

- Based on FIFO & Reference

- Steps:

- Look at the oldest page
  - Reference bit is 0 → remove it
  - Reference bit is 1
    - Change reference bit to 0
    - move to next oldest
- Repeat
- If all the frames referenced → change all references to 0 & remove the oldest one

Page Refs	3 Page Frames			
	Fault?	Page Contents		
A	yes	A*		
B	yes	B*	A*	
C	yes	C*	B*	A*
D	yes	D*	C	B
A	yes	A*	D*	C
B	yes	B*	A*	D*
E	yes	E*	B	A
A	no	E*	B	A*
B	no	E*	B*	A*
C	yes	C*	E	B
D	yes	D*	C*	E
E	no	D*	C*	E*

- **Exercise 1:** Program code occupies 1024 B in memory & uses a vector or 1000 B

- Physical memory: 1 MB
- Memory is paginated
- Page size: 512 B
- Addressing in 24 B
- Questions:

A. Indicate the

- 1) size of the virtual address space
- 2) offset
- 3) number of bits of the page number
- 4) number of bits of physical address
- 5) number of bits of physical frame
- 6) number of entries of PTE
- 7) size of page table

B. The loading of this page in memory causes internal fragmentation. Justify?

○ Answers:

A.

- 1) Maximum size of the virtual address space is  $2^{24}$
- 2)  $\text{Size}(\text{page}) = 512 \text{ B} = 2^9 \text{ B} \rightarrow$  I need 9 bits for the offset
- 3) 9 bits for offset  $\rightarrow$  (addressing- offset = page #)  $24-9 = 15$  bits for page number
- 4) Physical memory is  $1\text{MB} = 2^{20} \text{ B} \rightarrow$  20 bits for physical address
- 5) Bits for physical memory - bits for offset = bits for physical frame  $\rightarrow 20 - 9 = 11$  bits
- 6) number of PTE = number of pages =  $2^{15}$
- 7) size of page table = (number of PTE \* size of PTE) (size(PTE) is always 32 bits = 4B)  $\rightarrow 2^{15} * 2^2 = 2^{17}$

B. The process has code 1024 B & data 1000 B. The  $\text{size}(\text{frame}) = \text{size}(\text{page}) = 512 \text{ B}$

- 1) Code needs  $1024/512 = 2$  frames
- 2) Data needs 1 full frame (512 B) and some of another frame (488 B)
- 3)  $\rightarrow$  The process didn't use a full frame/page  $\rightarrow$  Internal fragmentation